

```

1 #%%
2 import pandas as pd
3 #%%
4 data_start = 0
5 data_end = 5000
6 df = pd.read_csv(
7         "/Users/krasimirtrifonov/Documents/GitHub/
8 Transformer-Neural-Network/dataset/
9 EURUSD_Daily_200005300000_202405300000.csv",
10        delimiter="\t")
11
12 # Extract the closing prices
13 closing = df["<CLOSE>"].iloc[data_start:data_end]
14 #%% md
15 ### Augmented Dickey-Fuller (ADF) test using the
16 statsmodels library:
17 #%% md
18 # Assume 'data' is your dataframe and '<CLOSE>' is
19 # the column of interest
20 #%% md
21 # Perform Augmented Dickey-Fuller test on the <
22 # CLOSE> column
23 adf_result = adfuller(df['<CLOSE>'])
24 # Extract the test statistics and p-value
25 adf_statistic = adf_result[0]
26 p_value = adf_result[1]
27 critical_values = adf_result[4]
28 # Print the results
29 print(f'ADF Statistic: {adf_statistic:.4f}')
30 print(f'p-value: {p_value:.4f}')
31 print('Critical Values:')
32 for key, value in critical_values.items():
33     print(f'    {key}: {value:.4f}')
34 #%% md
35         • The p-value of 0.253 is higher than the
36         usual significance levels (1%, 5%, or even 10%),
37         meaning we fail to reject the null hypothesis of

```

```

34 the test. This suggests that the data is non-
stationary.
35     • The ADF statistic is greater than all
critical values, further confirming that the series
likely has a unit root (non-stationary).
36 Let try to make it stationary by Differencing
and we perform the ADF test
37
38 #%%
39 # Differencing
40 close_diff = df['<CLOSE>'].diff().dropna()
41 #%% md
42 ### *Apply ADF on Difference Close
43 #%%
44 # Perform Augmented Dickey-Fuller test on the <
CLOSE> column
45 adf_result = adfuller(close_diff)
46
47 # Extract the test statistics and p-value
48 adf_statistic = adf_result[0]
49 p_value = adf_result[1]
50 critical_values = adf_result[4]
51
52 # Print the results
53 print(f'ADF Statistic on Difference Close: {'
      adf_statistic:.4f}')
```

54 print(f'p-value: {p_value:.4f}')

55 print('Critical Values:')

56 for key, value in critical_values.items():

57 print(f' {key}: {value:.4f}')

58 #%% md

59 • The ADF statistic is much lower than the
critical values, and the p-value is 0.0, meaning we
can now reject the null hypothesis.

60 • This indicates that the differenced data is
stationary.

61 #%% md

62 ### -Now let check for Engle's ARCH test to
check for **heteroscedasticity**

63 #%%

64 from statsmodels.stats.diagnostic import het_arch

```
65
66 # Perform Engle's ARCH test
67 arch_test_result = het_arch(close_diff)
68
69 # Extract the test statistic and p-value
70 arch_statistic = arch_test_result[0]
71 arch_p_value = arch_test_result[1]
72
73 # Print the results
74 print(f'ARCH Test Statistic on Close Diff: {arch_statistic:.4f}')
75 print(f'p-value: {arch_p_value:.4f}')
76 #%% md
77 ## • A low p-value (typically < 0.05) suggests heteroscedasticity (i.e., variance that changes over time).
78 ### Let try on repeating diff
79 #%%
80 close_diff_diff = close_diff.diff().dropna()
81 # Perform Engle's ARCH test on Repeating Close Diff
82 arch_test_result = het_arch(close_diff_diff)
83
84 # Extract the test statistic and p-value
85 arch_statistic = arch_test_result[0]
86 arch_p_value = arch_test_result[1]
87
88 # Print the results
89 print(f'ARCH Test Statistic on Repeating Close Diff: {arch_statistic:.4f}')
90 print(f'p-value: {arch_p_value:.9f}')
91
92 import matplotlib.pyplot as plt
93
94 # Plot the first 50 points of the second differenced data
95 plt.figure(figsize=(10, 6))
96
97 # Plot the second difference of <CLOSE>
98 plt.plot(close_diff_diff.head(50), marker='o', linestyle='--', label='Second Difference of <CLOSE>')
```

```
98 >')
99
100 # Add plot title and labels
101 plt.title('First 50 Points of Second Difference <
CLOSE>')
102 plt.xlabel('Index')
103 plt.ylabel('Second Difference Close')
104 plt.grid(True) # Adds a grid for better
    readability
105 plt.legend() # Shows the legend
106
107 # Display the plot
108 plt.show()
109 #%% md
110 ### • The extremely low p-value suggests that
    even after applying a second differencing, the
    data still exhibits heteroscedasticity (variance
    is not constant over time).
111 ## It is Time to apply GARCH
112 #%%
113 import pandas as pd
114 from arch import arch_model
115
116 # Load your data (ensure the percent change data
    is properly calculated)
117 # Assuming 'data' is the DataFrame with a '
    CLOSE_PERCENT_CHANGE' column
118
119 # Prepare the percent change data for GARCH
    modeling
120 garch_data = close_diff
121
122 # Fit a GARCH(1, 1) model on the percent change
    data
123 garch_model = arch_model(garch_data, vol='Garch',
    p=1, q=1)
124 garch_fitted = garch_model.fit(disp="off")
125
126 # Summarize the model
127 print(garch_fitted.summary())
128
```

```

129 # Forecast future volatility (for the next 5 days
   , for example)
130 forecast = garch_fitted.forecast(horizon=5)
131 print(forecast.variance[-1:]) # This gives the
   predicted variance (volatility) for the next 5
   periods
132 #%% md
133 The results of your GARCH(1,1) model are as
   follows, with some key interpretations and next
   steps:
134
135 GARCH Model Summary:
136
137 Mean Model:
138
139     • mu: The mean (constant) of the series is
   approximately 0.000024, with a p-value of 0.729,
   which indicates that the mean is not statistically
   significant.
140
141 Volatility (GARCH) Model:
142
143     • omega: 1.0101 \times 10^{-6} (p-value = 0.
   000) - This is the baseline level of variance, and
   it is statistically significant.
144     • alpha[1]: 0.0500 (p-value = 0.000) - This
   measures the influence of past shocks (squared
   returns or volatility) on future volatility. The
   significance indicates that past volatility
   influences future volatility.
145     • beta[1]: 0.9300 (p-value = 0.000) - This
   measures the persistence of volatility. A value
   close to 1 indicates high persistence, meaning
   that volatility tends to remain elevated for
   longer periods.
146
147 Key Observations:
148
149     1. High persistence of volatility: With a
   beta of 0.93, volatility is persistent, meaning
   that once volatility rises, it tends to remain

```

149 high for several periods.

150 2. Statistical significance: Both the alpha and beta coefficients are highly significant, meaning the GARCH model is capturing the volatility structure well.

151 3. Normal Distribution: The residuals (errors) are assumed to follow a normal distribution, but other distributions (like Student's t) might improve the fit.

152

153 Forecast:

154

155 The model produces forecasts for the next 5 periods of volatility:

156

157 • The forecasted volatility (variance) starts at 0.000021 and increases slightly to 0.000023 by the 5th period.

158

159 Convergence Warning:

160

161 • The message "Inequality constraints incompatible" suggests that the model did not converge fully. This could mean:

162 • The starting values or constraints may need adjustment.

163 • Trying different distributions (e.g., Student's t instead of normal) might help improve model fit.

164

165 Next Steps:

166

167 1. Refining the Model:

168 • Try changing the distribution assumption from normal to Student's t to allow for heavier tails.

169 • Ensure proper convergence by adjusting model parameters or re-running the optimization with different settings.

170 2. Volatility Interpretation:

171 • The forecasted volatility is low (around 0

```

171 .000021), indicating stable conditions in the
short term. If you're looking for predictive
insights, this information could be useful in
understanding future market behavior.
172     3. Residual Analysis:
173         • Check the residuals from the model to
ensure no significant patterns or autocorrelation
are left.
174 #%%
175 # Rescale the data (multiply by 100)
176 scaled_data = close_diff * 100
177
178 # Fit the GARCH model with rescaled data
179 garch_model = arch_model(scaled_data, vol='Garch'
, p=1, q=1)
180 garch_rescaled_fitted = garch_model.fit(disp="off"
)
181
182 # Summarize the model
183 print(garch_rescaled_fitted.summary())
184
185
186 # Forecasting future volatility with rescaled data
187 forecast = garch_rescaled_fitted.forecast(horizon=
5)
188 forecast_variance_rescaled = forecast.variance[-1
:] / 100 # Scale back
189 print(forecast_variance_rescaled)
190 #%% md
191 The recommendation to scale rather than normalize
might seem unusual because in many machine
learning contexts, normalization or
standardization is a common preprocessing step.
However, in the case of GARCH models (and other
econometric models), there are reasons why scaling
is preferred over normalization for better
optimization and interpretation:
192
193 Why Scaling is Recommended in GARCH Models:
194
195     1. Econometric Models Work with Raw Value

```

195 Ranges:

196 • In time series econometrics, especially with models like GARCH, the parameters are sensitive to the actual magnitude of the data. The variance and returns are calculated directly from the raw data, so the absolute scale of the input values affects how the optimizer estimates the volatility structure.

197 • Scaling by a constant factor (like multiplying by 100) maintains the relationships in the data, while simply normalizing (which changes the range to [0, 1]) may not preserve the underlying structure of volatility or returns in the same way.

198 2. Normalization Alters Statistical Properties:

199 • Normalizing the data (scaling it between 0 and 1 or z-scoring) would modify the mean and variance of the data. For models like GARCH, which rely on these properties to model volatility, such changes can disrupt the model's ability to accurately estimate parameters.

200 • Scaling by a constant factor (e.g., multiplying by 100) preserves the mean-to-variance ratio and leaves the overall structure of the time series unchanged.

201 3. Convergence Stability:

202 • Optimization algorithms in econometric models often perform better when the data is in a reasonable range (not too small or too large). If the data has very small values (as seen with your original data), the optimizer might struggle to find a solution due to the precision required for floating-point calculations. Rescaling moves the data into a more manageable range, improving convergence.

203 4. Maintaining Interpretability:

204 • When you scale the data by a constant (e.g., 100), the output from the GARCH model can still be interpreted in terms of the original scale. You can easily scale the results (e.g., volatility

```
204 forecasts) back to the original units.  
205     • In contrast, normalization would make  
interpreting the model's output more difficult, as  
it would no longer be in the original units of  
the data.  
206  
207 To Clarify:  
208  
209     • Scaling is a linear transformation that  
preserves the relationships between the data  
points (i.e., a proportional increase or decrease  
in values).  
210     • Normalization changes the distribution of  
the data, which can affect models that rely on the  
original statistical properties of the series (like  
variance, mean, and residuals).  
211  
212 Example:  
213  
214 If your data values are small (e.g.,  $5.05 \times 10^{-5}$ ), the optimizer might have trouble  
accurately estimating parameters because the  
changes between points are too subtle. Scaling by  
100 simply stretches the data but keeps the  
relationship between data points intact.  
215  
216 Summary:  
217  
218     • Normalization could alter the statistical  
properties (mean, variance) that are critical for  
GARCH modeling.  
219     • Scaling preserves the data's structure  
while making the optimization easier and improving  
convergence.  
220     • After fitting the model, you can easily  
scale the results back to the original units if  
needed.  
221 #%% md  
222 Step 1: Plot Conditional Volatility Over Time  
223  
224 You can plot the conditional variance (volatility
```

```
224 ) over time, which gives insight into how well the
      model captures changing volatility.
225 #%%
226 # Plot the conditional volatility over time (
      variance over time)
227 plt.figure(figsize=(10, 6))
228 # Plot the conditional volatility (square root of
      variance)
229 conditional_volatility = garch_rescaled_fitted.
      conditional_volatility
230 plt.plot(garch_rescaled_fitted.
      conditional_volatility, label='Conditional
      Volatility')
231 plt.title('Conditional Volatility Over Time (GARCH
      Model)')
232 plt.xlabel('Time')
233 plt.ylabel('Conditional Volatility')
234 plt.grid(True)
235 plt.legend()
236 plt.show()
237 #%% md
238 Step 2: Residual Analysis
239
240 After fitting the GARCH model, we can check
      whether the residuals are behaving like white
      noise (uncorrelated and normally distributed).
      This is a key test to verify if the model has
      captured the volatility dynamics adequately.
241 #%%
242 # Plot the standardized residuals
243 plt.figure(figsize=(10, 6))
244 resid = garch_rescaled_fitted.resid
245 plt.plot(garch_rescaled_fitted.resid, label='
      Standardized Residuals')
246 plt.title('Standardized Residuals from GARCH Model
      ')
247 plt.xlabel('Time')
248 plt.ylabel('Residuals')
249 plt.grid(True)
250 plt.legend()
251 plt.show()
```

```

252 #%%
253 from statsmodels.stats.diagnostic import
    acorr_ljungbox
254 import numpy as np
255 # Test for autocorrelation in the residuals using
    the Ljung-Box test
256 ljung_box_test = acorr_ljungbox(
    garch_rescaled_fitted.resid, lags=[10], return_df=
    True)
257
258 print(ljung_box_test)
259 if np.array(ljung_box_test.lb_pvalue)[0] > 0.05:
260     print('there's no significant autocorrelation
        in the residuals, suggesting that the GARCH model
        has adequately captured the volatility dynamics')
261 else:
262     print("GARCH is not adequate!")
263 %% md
264 If the p-values from the Ljung-Box test are
    greater than 0.05, it means there's no significant
    autocorrelation in the residuals, suggesting that
    the GARCH model has adequately captured the
    volatility dynamics.
265 %%
266 # Use a Q-Q plot to check if residuals are
    normally distributed
267 import scipy.stats as stats
268 stats.probplot(garch_rescaled_fitted.resid, dist="
    norm", plot=plt)
269 plt.title('Q-Q Plot of GARCH Model Residuals')
270 plt.show()
271 %% md
272 The Q-Q plot compares the distribution of the
    residuals to a theoretical normal distribution. In
    an ideal case where the residuals are normally
    distributed, the points should fall along the
    diagonal line.
273
274 In this case, deviations from the line, especially
    in the tails, suggest that the residuals might
    not be perfectly normally distributed. This could

```

```

274 indicate the presence of fat tails (kurtosis) or
skewness, which are common in financial time
series data.
275 #%%
276 # Forecast future volatility for 5 periods
277 forecast_rescaled = garch_rescaled_fitted.forecast
(horizon=5)
278
279 # Print forecasted variance
280 print("Forecasted variance over next 5 periods:")
281 print(forecast_rescaled.variance[-1:])
282 #%%
283 from scipy.stats import kurtosis, skew
284
285 # Assuming 'resid_data' contains the residuals
# from your GARCH model
286 # Replace 'resid_data['Residuals']' with your own
# data if necessary
287
288 # Calculate Skewness
289 skewness = skew(resid)
290
291 # Calculate Kurtosis
292 kurt = kurtosis(resid)
293
294 # Print the results
295 print(f'Skewness: {skewness:.4f}')
296 print(f'Kurtosis: {kurt:.4f}')
297 #%% md
298     Skewness: 0.0247
299     • A skewness value close to zero suggests
that the residuals are nearly symmetric. This
means there's no significant skewness in the data.
300     Kurtosis: 2.47
301     • Kurtosis less than 3 indicates platykurtic
behavior, meaning the residuals have lighter
tails than a normal distribution.
302     This suggests that extreme events (
outliers) are less frequent compared to what a
normal distribution would predict.
303     • The residuals appear to have little

```

```
303 skewness, meaning they are relatively symmetric
around the mean.
304     • The kurtosis is slightly lower than 3,
indicating fewer extreme values (fat tails)
compared to a normal distribution.
305 This suggests that the normal distribution
assumption might still hold, although it's on the
borderline.
306 #%% md
307 Step 1: Fit GARCH Models with Different
Distributions
308
309 You can fit the GARCH model with various
distributions to determine the best fit:
310 #%%
311 from arch import arch_model
312
313 # Assuming 'rescaled_data' contains your rescaled
# returns or percent changes
314 # Fit GARCH(1,1) with different distributions
315
316 # Normal distribution (default)
317 garch_normal = arch_model(scaled_data, vol='Garch',
, p=1, q=1, dist='normal').fit(disp="off")
318
319 # Student's t-distribution
320 garch_t = arch_model(scaled_data, vol='Garch', p=1
, q=1, dist='t').fit(disp="off")
321
322 # Skewed Student's t-distribution
323 garch_skewt = arch_model(scaled_data, vol='Garch'
, p=1, q=1, dist='skewt').fit(disp="off")
324
325 # Print model summaries
326 print("Normal Distribution GARCH:")
327 print(garch_normal.summary())
328
329 print("\nStudent's t-Distribution GARCH:")
330 print(garch_t.summary())
331
332 print("\nSkewed t-Distribution GARCH:")
```

```

333 print(garch_skewt.summary())
334 #%% md
335 Step 2: Compare the Distributions Using AIC
336
337 The AIC (Akaike Information Criterion) can be used
      to compare the different models. The model with
      the lowest AIC is considered the best fit.
338 #%%
339 # Compare AIC values for the different models
340 print(f"Normal AIC: {garch_normal.aic}")
341 print(f"Student's t AIC: {garch_t.aic}")
342 print(f"Skewed t AIC: {garch_skewt.aic}")
343 #%% md
344     • The Student's t-distribution is likely
       capturing the fat tails better, which means that
       your data contains
345     more extreme values than a normal
       distribution would predict.
346     Since the Student's t-distribution GARCH
       model has the best fit, let's move forward with a
       deeper analysis. Here's a breakdown of what we can
       do next:
347
348 1. Residual Analysis for the Student's t-
       Distribution GARCH Model:
349
350     • We need to ensure that the residuals from
       the Student's t-distribution GARCH model behave
       like white noise (i.e., no autocorrelation and are
       approximately normally distributed).
351
352 2. Forecast Volatility Using the Student's t-
       Distribution GARCH Model:
353
354     • We can also forecast future volatility
       using this model to see how well it captures
       future uncertainty.
355 #%%
356 # Plot the standardized residuals from the Student
      's t-distribution GARCH model
357 plt.figure(figsize=(10, 6))

```

```
358 plt.plot(garch_t.resid, label='Standardized  
Residuals (Student\'s t GARCH)')  
359 plt.title('Standardized Residuals from Student\'s  
t GARCH Model')  
360 plt.xlabel('Time')  
361 plt.ylabel('Residuals')  
362 plt.grid(True)  
363 plt.legend()  
364 plt.show()  
365 #%% md  
366 Step 2: Test for Autocorrelation in the Residuals  
367  
368 We'll use the Ljung-Box test to check if there is  
any remaining autocorrelation in the residuals.  
Ideally, there should be no significant  
autocorrelation left after the GARCH model is  
applied.  
369 #%%  
370
```